

Self-Driving Cars


Ex. 01 - Coding Challenge - Imitation Learning

Katrin Renz

Autonomous Vision Group

University of Tübingen

EBERHARD KARLS
UNIVERSITÄT
TÜBINGEN



e l i a s
European Laboratory for Learning and Intelligent Systems

Questions?

Exercise Setup

Files

- ▶ `01_imitation_learning_exercise.pdf`
- ▶ `main.py`
- ▶ `network.py`, `training.py`, `demonstrations.py`
- ▶ `submission.txt`
- ▶ `data`, a folder with demonstrations

Exercise Setup

Files

- ▶ `01_imitation_learning_exercise.pdf`
- ▶ `main.py`
- ▶ `network.py`, `training.py`, `demonstrations.py`
- ▶ `submission.txt`
- ▶ `data`, a folder with demonstrations

Submit

- ▶ submission information: please fill out `submission.txt`
- ▶ your code: `network.py`, `training.py`, `demonstrations.py` as a `.zip` file
- ▶ your pre-trained model as a `.t7` file

Deadline: **Tue, 16. November 2021 - 8pm**

Exercise Setup

Do's

- ▶ comment your code
- ▶ use docstrings
- ▶ use self-explanatory variable names
- ▶ structure your code well

Exercise Setup

Do's

- ▶ comment your code
- ▶ use docstrings
- ▶ use self-explanatory variable names
- ▶ structure your code well

Do not's

- ▶ change `main.py`, especially `calculate_score_for_leaderboard()`
- ▶ install more packages
- ▶ change the gym environment

Imitation Learning

Behavioral Cloning

Imitation Learning

Components:

- ▶ State: $s \in \mathcal{S}$ may be partially observed (e.g., game screen)
- ▶ Action: $a \in \mathcal{A}$ may be discrete or continuous (e.g., turn angle, speed)
- ▶ Policy: $\pi_\theta : \mathcal{S} \rightarrow \mathcal{A}$ we want to learn the policy parameters θ
- ▶ Optimal action: $a^* \in \mathcal{A}$ provided by expert demonstrator
- ▶ Optimal policy: $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ provided by expert demonstrator
- ▶ State dynamics: $P(s_{i+1}|s_i, a_i)$ simulator, typically not known to policy
Often deterministic: $s_{i+1} = T(s_i, a_i)$ deterministic mapping
- ▶ Rollout: Given s_0 , sequentially execute $a_i = \pi_\theta(s_i)$ & sample $s_{i+1} \sim P(s_{i+1}|s_i, a_i)$
yields trajectory $\tau = (s_0, a_0, s_1, a_1, \dots)$
- ▶ Loss function: $\mathcal{L}(a^*, a)$ loss of action a given optimal action a^*

Imitation Learning

Components:

- ▶ State: $s \in \mathcal{S}$ may be partially observed (e.g., game screen)
- ▶ Action: $a \in \mathcal{A}$ may be discrete or continuous (e.g., turn angle, speed)
- ▶ Policy: $\pi_\theta : \mathcal{S} \rightarrow \mathcal{A}$ we want to learn the policy parameters θ
- ▶ Optimal action: $a^* \in \mathcal{A}$ provided by expert demonstrator
- ▶ Optimal policy: $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ provided by expert demonstrator
- ▶ State dynamics: $P(s_{i+1}|s_i, a_i)$ simulator, typically not known to policy
Often deterministic: $s_{i+1} = T(s_i, a_i)$ deterministic mapping
- ▶ Rollout: Given s_0 , sequentially execute $a_i = \pi_\theta(s_i)$ & sample $s_{i+1} \sim P(s_{i+1}|s_i, a_i)$
yields trajectory $\tau = (s_0, a_0, s_1, a_1, \dots)$
- ▶ Loss function: $\mathcal{L}(a^*, a)$ loss of action a given optimal action a^*

1.1

Network Design

a) Load demonstrations

```
def load_demonstrations(data_folder):  
    """  
    1.1 a)  
    Given the folder containing the expert demonstrations, the data gets loaded and  
    stored it in two lists: observations and actions.  
    N = number of (observation, action) - pairs  
    data_folder:    python string, the path to the folder containing the  
                    observation_%05d.npy and action_%05d.npy files  
    return:  
    observations:    python list of N numpy.ndarrays of size (96, 96, 3)  
    actions:         python list of N numpy.ndarrays of size 3  
    """  
    pass
```


b) Understand training

```
def train(data_folder, trained_network_file):  
    """  
    Function for training the network.  
    """  
    infer_action = ClassificationNetwork()  
    optimizer = torch.optim.Adam(infer_action.parameters(), lr=1e-2)  
    loss_function = nn.CrossEntropyLoss()  
  
    observations, actions = load_demonstrations(data_folder)  
    observations = [torch.Tensor(observation) for observation in observations]  
    actions = [torch.Tensor(action) for action in actions]  
  
    batches = [batch for batch in zip(observations,  
                                     infer_action.actions_to_classes(actions))]  
  
    # setting device on GPU if available, else CPU  
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
  
    nr_epochs = 50  
    batch_size = 64  
    start_time = time.time()  
  
    for epoch in range(nr_epochs):  
        random.shuffle(batches)
```

c) Classification Network

- ▶ expert demonstrations
action = [steer, gas, brake]
e.g. [1., 0., 0.8]
- ▶ define **action-classes**
 - ▶ {steer_left}
 - ▶ {}
 - ▶ {steer_right, brake}
 - ▶ {gas}

⇒ map [1., 0., 0.8] → ?



c) Classification Network

- ▶ `actions_to_classes`
expert demonstrations \rightarrow action-classes
- ▶ `scores_to_action`
score predicted by the network \rightarrow action
- ▶ `CrossEntropyLoss`
loss function: gt vs. prediction



d) Implement network

```
class ClassificationNetwork(torch.nn.Module):
    def __init__(self):
        """
        1.1 d)
        Implementation of the network layers. The image size of the input
        observations is 96x96 pixels.
        """
        super().__init__()

        # setting device on GPU if available, else CPU
        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

    def forward(self, observation):
        """
        1.1 e)
        The forward pass of the network. Returns the prediction for the given
        input observation.
        observation: torch.Tensor of size (batch_size, 96, 96, 3)
        return      torch.Tensor of size (batch_size, C)
        """
        pass
```

- ▶ 2 to 3 convolution layers + ReLU
- ▶ 2 to 3 fully connected layers + ReLU

e) Forward pass, train and test

- ▶ color channels or gray-scale
- ▶ `python main.py --train`
- ▶ `python main.py --test`
- ▶ hyper-parameter tuning

e) Forward pass, train and test

python main.py --test


```
def evaluate(trained_network_file):  
    """  
    """  
    infer_action = torch.load(trained_network_file)  
    infer_action.eval()  
    env = gym.make('CarRacing-v0')  
  
    # setting device on GPU if available, else CPU  
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
  
    infer_action = infer_action.to(device)  
  
    for episode in range(5):  
        observation = env.reset()  
  
        reward_per_episode = 0  
        for t in range(500):  
            env.render()  
            action_scores = infer_action(torch.Tensor(  
                np.ascontiguousarray(observation[None])).to(device))  
  
            steer, gas, brake = infer_action.scores_to_action(action_scores)  
            observation, reward, done, info = env.step([steer, gas, brake])  
            reward_per_episode += reward  
  
        print('episode %d \t reward %f' % (episode, reward_per_episode))
```

f) Record own demonstrations


```
def record_demonstrations(demonstrations_folder):  
    """  
    Function to record own demonstrations by driving the car in the gym car-racing  
    environment.  
    demonstrations_folder: python string, the path to where the recorded demonstrations  
    | | | | | are to be saved  
  
    The controls are:  
    arrow keys: control the car; steer left, steer right, gas, brake  
    ESC: quit and close  
    SPACE: restart on a new track  
    TAB: save the current run  
    """  
    env = gym.make('CarRacing-v0').env  
    status = ControlStatus()
```

f) Record own demonstrations

```
python main.py --teach
```



```
python main.py --test
```



1.2

Network Improvements

Competition

Competition

python main.py --score

```
def calculate_score_for_leaderboard(trained_network_file):  
    """  
    Evaluate the performance of the network. This is the function to be used for  
    the final ranking on the course-wide leader-board, only with a different set  
    of seeds. Better not change it.  
    """  
  
    infer_action = torch.load(trained_network_file)  
    infer_action.eval()  
    env = gym.make('CarRacing-v0')  
  
    # setting device on GPU if available, else CPU  
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
  
    seeds = [22597174, 68545857, 75568192, 91140053, 86018367,  
            49636746, 66759182, 91294619, 84274995, 31531469]  
  
    total_reward = 0  
  
    for episode in range(10):  
        env.seed(seeds[episode])  
        observation = env.reset()
```

Questions?